

Probability Practice :

Part A :

Visitors to your website are asked to answer a single survey question before they get access to the content on the page. Among all of the users, there are two categories: Random Clicker (RC), and Truthful Clicker (TC). There are two possible answers to the survey: yes and no. Random clickers would click either one with equal probability. You are also giving the information that the expected fraction of random clickers is 0.3.

After a trial period, you get the following survey results: 65% said Yes and 35% said No.

What fraction of people who are truthful clickers answered yes?

Probability of yes for a random clicker. $P(Y/R) = 0.5$

Fraction of random clicker $P(R) = 0.3$

Fraction of truthful clicker $P(T) = 0.7$

Fraction of yes $P(Y) = 0.65$

Fraction of no $P(N) = 0.35$

Fraction of yes explained as a sum of conditional probability $P(Y) = P(Y/R) \times P(R) + P(Y/T) \times P(T)$

$$0.65 = 0.5 \times 0.3 + P(Y/T) \times 0.7$$

Fraction of truthful clickers that answered yes $P(Y/T) = 0.714$

Part B :

Imagine a medical test for a disease with the following two attributes:

The sensitivity is about 0.993. That is, if someone has the disease, there is a probability of 0.993 that they will test positive. The specificity is about 0.9999. This means that if someone doesn't have the disease, there is probability of 0.9999 that they will test negative. In the general population, incidence of the disease is reasonably rare: about 0.0025% of all people have it (or 0.000025 as a decimal probability).

Suppose someone tests positive. What is the probability that they have the disease? In light of this calculation, do you envision any problems in implementing a universal testing policy for the disease?

$P(D)$ = Probability of having the disease

$P(T)$ = Probability of testing positive.

$P(N)$ = Probability of testing negative

$P(ND)$ = Probability of not having disease

$P(T/D)$ = Probability of testing positive given that they have the disease

$P(N/ND)$ = Probability of testing negative given that they do not have the disease

$$P(D) = 0.000025$$

$$P(T/D) = 0.993$$

$$P(N/ND) = 0.9999$$

We use Bayes theorem to calculate the probability, $P(D/T)$ which is probability of having the disease given that the test is positive.

$$P(D/T) = P(T/D) \times P(D) / P(T)$$

$$P(T) = P(T/D) \times P(D) + (1 - P(N/ND)) \times (1 - P(D))$$

$$P(D/T) = 0.993 \times 0.000025 / (0.993 \times 0.000025 + (1 - 0.9999) \times (1 - 0.000025))$$

$$P(D/T) = 0.1988$$

If a universal testing policy is implemented, then the chance that they actually have the disease when they are tested positive is 0.198 which is a low value. The probability that the person does not have the disease is very high ($1 - 0.000025$). Because of this, if a user tests positive, then it is more likely that they do not have the disease than they do.

Green buildings

```
library(ggplot2)
library(lattice)
library(mosaic)

## Loading required package: dplyr

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

## Loading required package: mosaicData

## Loading required package: Matrix

##
## The 'mosaic' package masks several functions from core packages in order to
## add additional features.
## The original behavior of these functions should not be affected by this.
##
## Attaching package: 'mosaic'

## The following object is masked from 'package:Matrix':
##
##   mean
```

```

## The following objects are masked from 'package:dplyr':
##
##   count, do, tally

## The following objects are masked from 'package:stats':
##
##   binom.test, cor, cov, D, fivenum, IQR, median, prop.test,
##   quantile, sd, t.test, var

## The following objects are masked from 'package:base':
##
##   max, mean, min, prod, range, sample, sum

library(RCurl)

## Loading required package: bitops

set.seed(100)
rm(list=ls()) #Clear Workspace
temp = getURL("https://raw.githubusercontent.com/matt-staton/stat_380/master/greenbuildings.csv")
greenbuildings = read.csv(text = temp, header=T)
gbuild = greenbuildings

attach(gbuild)
gbuild$Rent_Diff = Rent - cluster_rent
names(gbuild)

## [1] "CS_PropertyID"      "cluster"            "size"
## [4] "empl_gr"           "Rent"               "leasing_rate"
## [7] "stories"           "age"                "renovated"
## [10] "class_a"           "class_b"            "LEED"
## [13] "Energystar"        "green_rating"        "net"
## [16] "amenities"         "cd_total_07"         "hd_total07"
## [19] "total_dd_07"       "Precipitation"      "Gas_Costs"
## [22] "Electricity_Costs" "cluster_rent"       "Rent_Diff"

summary(gbuild)

##   CS_PropertyID      cluster      size      empl_gr
##   Min.   :      1   Min.   :  1.0   Min.   : 1624   Min.   : -24.950
##   1st Qu.: 157452   1st Qu.: 272.0   1st Qu.: 50891   1st Qu.:  1.740
##   Median : 313253   Median : 476.0   Median : 128838   Median :  1.970
##   Mean    : 453003   Mean    : 588.6   Mean    : 234638   Mean    :  3.207
##   3rd Qu.: 441188   3rd Qu.:1044.0   3rd Qu.: 294212   3rd Qu.:  2.380
##   Max.    :6208103   Max.    :1230.0   Max.    :3781045   Max.    : 67.780
##                                     NA's    :74
##           Rent      leasing_rate      stories      age
##   Min.   :  2.98   Min.   :  0.00   Min.   :  1.00   Min.   :  0.00
##   1st Qu.: 19.50   1st Qu.: 77.85   1st Qu.:  4.00   1st Qu.: 23.00
##   Median : 25.16   Median : 89.53   Median : 10.00   Median : 34.00
##   Mean    : 28.42   Mean    : 82.61   Mean    : 13.58   Mean    : 47.24

```

```
## 3rd Qu.: 34.18 3rd Qu.: 96.44 3rd Qu.: 19.00 3rd Qu.: 79.00
## Max. :250.00 Max. :100.00 Max. :110.00 Max. :187.00
##
## renovated class_a class_b LEED
## Min. :0.0000 Min. :0.0000 Min. :0.0000 Min. :0.000000
## 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.0000 1st Qu.:0.000000
## Median :0.0000 Median :0.0000 Median :0.0000 Median :0.000000
## Mean :0.3795 Mean :0.3999 Mean :0.4595 Mean :0.006841
## 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:1.0000 3rd Qu.:0.000000
## Max. :1.0000 Max. :1.0000 Max. :1.0000 Max. :1.000000
##
## Energystar green_rating net amenities
## Min. :0.00000 Min. :0.00000 Min. :0.00000 Min. :0.0000
## 1st Qu.:0.00000 1st Qu.:0.00000 1st Qu.:0.00000 1st Qu.:0.0000
## Median :0.00000 Median :0.00000 Median :0.00000 Median :1.0000
## Mean :0.08082 Mean :0.08677 Mean :0.03471 Mean :0.5266
## 3rd Qu.:0.00000 3rd Qu.:0.00000 3rd Qu.:0.00000 3rd Qu.:1.0000
## Max. :1.00000 Max. :1.00000 Max. :1.00000 Max. :1.0000
##
## cd_total_07 hd_total07 total_dd_07 Precipitation
## Min. : 39 Min. : 0 Min. :2103 Min. :10.46
## 1st Qu.: 684 1st Qu.:1419 1st Qu.:2869 1st Qu.:22.71
## Median : 966 Median :2739 Median :4979 Median :23.16
## Mean :1229 Mean :3432 Mean :4661 Mean :31.08
## 3rd Qu.:1620 3rd Qu.:4796 3rd Qu.:6413 3rd Qu.:43.89
## Max. :5240 Max. :7200 Max. :8244 Max. :58.02
##
## Gas_Costs Electricity_Costs cluster_rent Rent_Diff
## Min. :0.009487 Min. :0.01780 Min. : 9.00 Min. : -45.9150
## 1st Qu.:0.010296 1st Qu.:0.02330 1st Qu.:20.00 1st Qu.: -2.9650
## Median :0.010296 Median :0.03274 Median :25.14 Median : 0.0000
## Mean :0.011336 Mean :0.03096 Mean :27.50 Mean : 0.9213
## 3rd Qu.:0.011816 3rd Qu.:0.03781 3rd Qu.:34.00 3rd Qu.: 3.2800
## Max. :0.028914 Max. :0.06280 Max. :71.44 Max. :191.2800
##
```

```
lm.fit = lm(Rent ~., data = gbuild)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = Rent ~ ., data = gbuild)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -5.284e-12 -3.500e-15  2.000e-16  3.200e-15  5.451e-12
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  7.172e-14  1.136e-14  6.311e+00 2.92e-10 ***
```

```

## CS_PropertyID      1.293e-20  1.750e-21  7.387e+00  1.65e-13 ***
## cluster            1.643e-17  3.157e-18  5.204e+00  2.00e-07 ***
## size              -1.023e-19  7.341e-21 -1.393e+01  < 2e-16 ***
## empl_gr           2.160e-15  1.891e-16  1.142e+01  < 2e-16 ***
## leasing_rate      5.162e-16  5.927e-17  8.711e+00  < 2e-16 ***
## stories           -1.470e-16  1.798e-16 -8.180e-01  0.4136
## age               4.262e-16  5.245e-17  8.126e+00  5.14e-16 ***
## renovated        -2.132e-14  2.874e-15 -7.418e+00  1.31e-13 ***
## class_a           3.958e-14  4.878e-15  8.116e+00  5.57e-16 ***
## class_b           6.067e-15  3.812e-15  1.592e+00  0.1115
## LEED              1.309e-14  3.981e-14  3.290e-01  0.7422
## Energystar        5.628e-14  4.243e-14  1.326e+00  0.1847
## green_rating     -7.069e-14  4.266e-14 -1.657e+00  0.0975 .
## net              -2.460e-15  6.597e-15 -3.730e-01  0.7093
## amenities        -4.580e-15  2.801e-15 -1.636e+00  0.1020
## cd_total_07      -1.284e-17  1.628e-18 -7.890e+00  3.42e-15 ***
## hd_total07        3.904e-19  9.994e-19  3.910e-01  0.6961
## total_dd_07             NA             NA             NA             NA
## Precipitation     2.724e-15  1.792e-16  1.520e+01  < 2e-16 ***
## Gas_Costs        -9.165e-12  8.727e-13 -1.050e+01  < 2e-16 ***
## Electricity_Costs 5.655e-12  2.781e-13  2.033e+01  < 2e-16 ***
## cluster_rent      1.000e+00  1.580e-16  6.331e+15  < 2e-16 ***
## Rent_Diff         1.000e+00  1.259e-16  7.946e+15  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.046e-13 on 7797 degrees of freedom
## (74 observations deleted due to missingness)
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 7.408e+30 on 22 and 7797 DF,  p-value: < 2.2e-16

```

confint(lm.fit)

```

##              2.5 %          97.5 %
## (Intercept)  4.944170e-14  9.399177e-14
## CS_PropertyID 9.497826e-21  1.635919e-20
## cluster      1.023998e-17  2.261785e-17
## size        -1.166696e-19 -8.789077e-20
## empl_gr      1.789213e-15  2.530539e-15
## leasing_rate 4.000657e-16  6.324173e-16
## stories     -4.993372e-16  2.053959e-16
## age         3.233495e-16  5.289668e-16
## renovated   -2.695144e-14 -1.568444e-14
## class_a     3.002321e-14  4.914599e-14
## class_b    -1.405133e-15  1.353961e-14
## LEED        -6.494246e-14  9.113091e-14
## Energystar  -2.689349e-14  1.394475e-13
## green_rating -1.543178e-13  1.293406e-14
## net         -1.539271e-14  1.047289e-14
## amenities   -1.007010e-14  9.093739e-16

```

```
## cd_total_07      -1.603215e-17 -9.651347e-18
## hd_total07       -1.568757e-18  2.349594e-18
## total_dd_07      NA              NA
## Precipitation    2.372503e-15  3.075061e-15
## Gas_Costs        -1.087532e-11 -7.453807e-12
## Electricity_Costs 5.109408e-12  6.199704e-12
## cluster_rent     1.000000e+00  1.000000e+00
## Rent_Diff        1.000000e+00  1.000000e+00
```

We can see here controlling for all available variables that the most significant predictors of price are PropertyID, cluster (ie: location), size, employment growth, cluster Rent, stories, leasing rate, hd-total-07 (total heating days in 2007), precipitation, Gas costs, age, class A, class B, net, amenities and electricity costs.

These predictors seem intuitive with the exception of green_rating having almost no predictive power. The green rating has a 95% CI of -6.827645e+00 8.221535e+00, which includes 0, leaving us to accept the null hypothesis that green rating is not statistically significant. Furthermore, I was very surprised to see renovation having very little predictive power, with a 95% confidence interval of -6.493550e-01 3.644221e-01.

```
names(lm.fit)
```

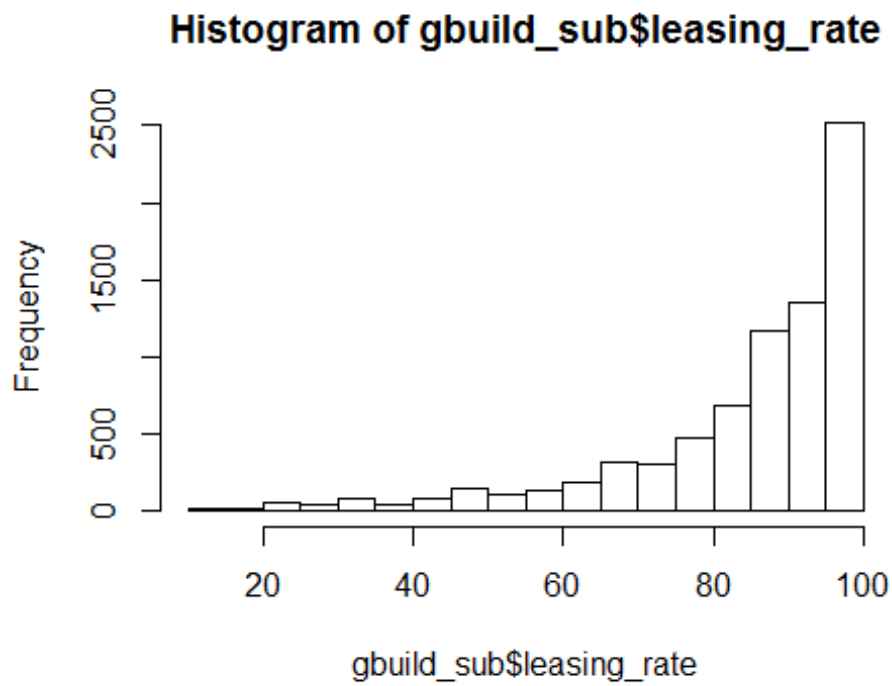
```
## [1] "coefficients" "residuals"      "effects"        "rank"
## [5] "fitted.values" "assign"          "qr"             "df.residual"
## [9] "na.action"     "xlevels"         "call"           "terms"
## [13] "model"
```

```
hist(gbuild$leasing_rate, plot=TRUE)
```



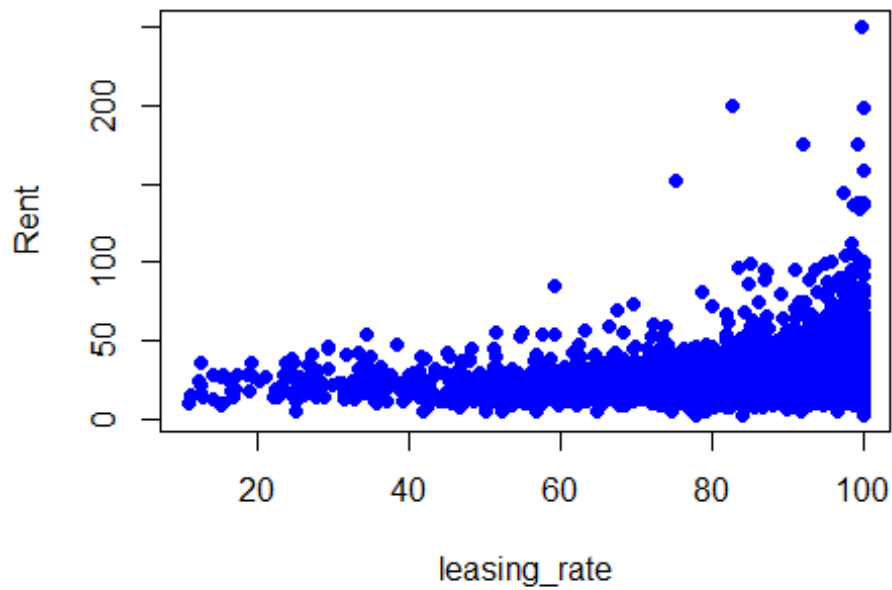
In order to do an apples to apples comparison of the previous analysis I will re-run the model with leasing rates $\geq 10\%$. Let's also plot rent as function of leasing rate to understand the effect removing the bottom 10 percentil will have.

```
gbuild_sub = subset(gbuild, gbuild$leasing_rate >= 10) #Remove Lease rates < 10%  
  
hist(gbuild_sub$leasing_rate, plot=TRUE)
```



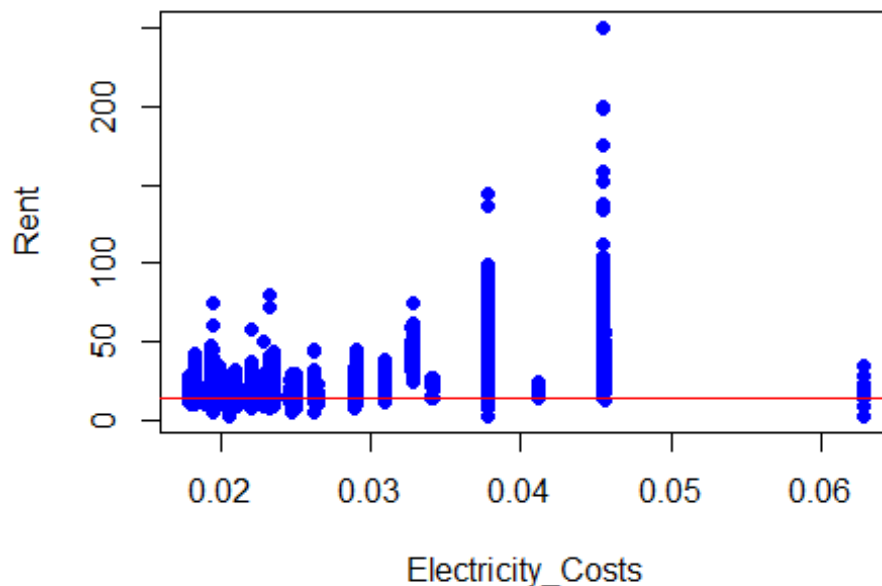
#Rent positively correlated with Leasing rate

```
plot(Rent~leasing_rate,data = gbuild_sub, col="blue",pch=16)
```




```
#Rent positively correlated with electricity costs
plot(Rent~Electricity_Costs,data = gbuild_sub, col="blue",pch=16)

abline(lm(Rent~leasing_rate,data = gbuild_sub),col="red")
```



```
gbuild_sub$util_index = gbuild_sub$hd_total07*gbuild_sub$Gas_Costs +
  gbuild_sub$cd_total_07 * gbuild_sub$Electricity_Costs
lm.fit2 = lm(Rent ~.+util_index*class_a+cd_total_07*class_a+green_rating*clas
s_a+empl_gr*class_a+empl_gr*green_rating, data = gbuild_sub)
summary(lm.fit2)
```

```
##
## Call:
## lm(formula = Rent ~ . + util_index * class_a + cd_total_07 *
##   class_a + green_rating * class_a + empl_gr * class_a + empl_gr *
##   green_rating, data = gbuild_sub)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.221e-12 -1.100e-14  1.000e-15  1.040e-14  2.559e-11
##
## Coefficients: (1 not defined because of singularities)
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   1.696e-13  7.403e-14  2.291e+00  0.02202 *
## CS_PropertyID  2.623e-20  5.412e-21  4.847e+00  1.28e-06 ***
## cluster       -2.619e-17  9.068e-18 -2.889e+00  0.00388 **
## size          -3.519e-20  2.088e-20 -1.685e+00  0.09197 .
```

```

## empl_gr -3.503e-16 7.524e-16 -4.660e-01 0.64158
## leasing_rate 3.802e-16 2.150e-16 1.768e+00 0.07704 .
## stories 2.598e-16 5.126e-16 5.070e-01 0.61235
## age 8.216e-17 1.540e-16 5.330e-01 0.59378
## renovated 1.577e-16 8.302e-15 1.900e-02 0.98484
## class_a -1.894e-14 2.036e-14 -9.310e-01 0.35213
## class_b -6.154e-15 1.116e-14 -5.510e-01 0.58147
## LEED 2.406e-14 1.129e-13 2.130e-01 0.83118
## Energystar 5.744e-14 1.203e-13 4.770e-01 0.63309
## green_rating -3.586e-14 1.223e-13 -2.930e-01 0.76940
## net -2.120e-14 1.886e-14 -1.124e+00 0.26101
## amenities 7.752e-15 8.077e-15 9.600e-01 0.33719
## cd_total_07 2.097e-17 1.892e-17 1.108e+00 0.26774
## hd_total07 7.324e-18 7.790e-18 9.400e-01 0.34718
## total_dd_07 NA NA NA NA
## Precipitation -5.308e-16 5.195e-16 -1.022e+00 0.30693
## Gas_Costs 5.085e-12 4.338e-12 1.172e+00 0.24116
## Electricity_Costs -5.203e-13 1.156e-12 -4.500e-01 0.65278
## cluster_rent 1.000e+00 4.620e-16 2.164e+15 < 2e-16 ***
## Rent_Diff 1.000e+00 3.592e-16 2.784e+15 < 2e-16 ***
## util_index -5.825e-16 5.953e-16 -9.790e-01 0.32781
## class_a:util_index -6.081e-16 2.966e-16 -2.050e+00 0.04039 *
## class_a:cd_total_07 5.131e-17 1.184e-17 4.332e+00 1.50e-05 ***
## class_a:green_rating 3.311e-14 2.966e-14 1.116e+00 0.26427
## empl_gr:class_a -2.316e-15 1.017e-15 -2.278e+00 0.02275 *
## empl_gr:green_rating -5.280e-16 1.330e-15 -3.970e-01 0.69150
## ---
## Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.963e-13 on 7577 degrees of freedom
## (73 observations deleted due to missingness)
## Multiple R-squared: 1, Adjusted R-squared: 1
## F-statistic: 7.118e+29 on 28 and 7577 DF, p-value: < 2.2e-16

```

confint(lm.fit2)

```

##          2.5 %          97.5 %
## (Intercept) 2.444902e-14 3.146911e-13
## CS_PropertyID 1.562321e-20 3.683954e-20
## cluster -4.396784e-17 -8.417897e-18
## size -7.613067e-20 5.742406e-21
## empl_gr -1.825168e-15 1.124666e-15
## leasing_rate -4.125658e-17 8.016243e-16
## stories -7.450957e-16 1.264617e-15
## age -2.197758e-16 3.840877e-16
## renovated -1.611641e-14 1.643183e-14
## class_a -5.884580e-14 2.096202e-14
## class_b -2.803693e-14 1.572905e-14
## LEED -1.971601e-13 2.452805e-13
## Energystar -1.784055e-13 2.932774e-13

```

```
## green_rating      -2.756190e-13  2.039027e-13
## net               -5.818207e-14  1.577347e-14
## amenities         -8.080561e-15  2.358441e-14
## cd_total_07       -1.612036e-17  5.806592e-17
## hd_total07        -7.947044e-18  2.259476e-17
## total_dd_07              NA          NA
## Precipitation     -1.549050e-15  4.875257e-16
## Gas_Costs         -3.418602e-12  1.358805e-11
## Electricity_Costs -2.787123e-12  1.746545e-12
## cluster_rent      1.000000e+00  1.000000e+00
## Rent_Diff         1.000000e+00  1.000000e+00
## util_index        -1.749422e-15  5.843577e-16
## class_a:util_index -1.189638e-15 -2.664067e-17
## class_a:cd_total_07 2.808932e-17  7.452550e-17
## class_a:green_rating -2.502430e-14  9.124126e-14
## empl_gr:class_a     -4.308981e-15 -3.231612e-16
## empl_gr:green_rating -3.135996e-15  2.080056e-15
```

We can see leasing rate has a positive effect on rent prices, we won't do anything with that for now, but we may need to control for that in the future.

Again we see none of the green IV's have any reliable predictive power as it relates to price. The most meaningful predictors remain to be Size, Employment growth, Class A, Class B, Net, HD-total07, Gas-Costs, Electricity-Costs, and Cluster_Rent (IE: Neighborhood average rent)

The analysis doesn't account for lurking variables by just looking at median rent. It assumes the higher prices of green buildings are due to them being green, when we can clearly see from the regression model, they are not. This is because the regression model shows the marginal effect of each variable, and allows one to control for other factors. Furthermore, in order to reliably predict the price of the building, the developer should input the values of the most important predictors above to estimate. Using the median however was a good idea, because it is more robust to outliers.

Lets remove all statistically insignificant variables using step-wise regression; acknowledging that the coefficients may change slightly

```
lm.fit3 = step(lm.fit2 , scope=formula(lm.fit2), direction="back", k=log(length(gbuild_sub)))
```

```
## Start:  AIC=-438764.6
## Rent ~ CS_PropertyID + cluster + size + empl_gr + leasing_rate +
##      stories + age + renovated + class_a + class_b + LEED + Energystar +
##      green_rating + net + amenities + cd_total_07 + hd_total07 +
##      total_dd_07 + Precipitation + Gas_Costs + Electricity_Costs +
##      cluster_rent + Rent_Diff + util_index + util_index * class_a +
##      cd_total_07 * class_a + green_rating * class_a + empl_gr *
##      class_a + empl_gr * green_rating
```

```

## Warning: attempting model selection on an essentially perfect fit is
## nonsense

##
## Step: AIC=-438764.6
## Rent ~ CS_PropertyID + cluster + size + empl_gr + leasing_rate +
##   stories + age + renovated + class_a + class_b + LEED + Energystar +
##   green_rating + net + amenities + cd_total_07 + hd_total07 +
##   Precipitation + Gas_Costs + Electricity_Costs + cluster_rent +
##   Rent_Diff + util_index + class_a:util_index + class_a:cd_total_07 +
##   class_a:green_rating + empl_gr:class_a + empl_gr:green_rating

## Warning: attempting model selection on an essentially perfect fit is
## nonsense

##           Df Sum of Sq    RSS    AIC
## - empl_gr:green_rating  1         0      0 -438767
## - class_a:green_rating  1         0      0 -438766
## <none>                                0 -438765
## - class_a:util_index    1         0      0 -438763
## - empl_gr:class_a       1         0      0 -438762
## - LEED                  1         0      0 -438760
## - CS_PropertyID         1         0      0 -438753
## - class_a:cd_total_07   1         0      0 -438749
## - Electricity_Costs     1         0      0 -438737
## - leasing_rate          1         0      0 -438736
## - renovated             1         0      0 -438723
## - Energystar            1         0      0 -438715
## - hd_total07            1         0      0 -438714
## - stories               1         0      0 -438711
## - Precipitation         1         0      0 -438655
## - class_b               1         0      0 -438629
## - net                   1         0      0 -438586
## - amenities             1         0      0 -438520
## - Gas_Costs             1         0      0 -438519
## - age                   1         0      0 -438514
## - size                  1         0      0 -438466
## - cluster               1         0      0 -438168
## - cluster_rent          1    411137 411137   30438
## - Rent_Diff             1    680136 680136   34267
##
## Step: AIC=-438767.6
## Rent ~ CS_PropertyID + cluster + size + empl_gr + leasing_rate +
##   stories + age + renovated + class_a + class_b + LEED + Energystar +
##   green_rating + net + amenities + cd_total_07 + hd_total07 +
##   Precipitation + Gas_Costs + Electricity_Costs + cluster_rent +
##   Rent_Diff + util_index + class_a:util_index + class_a:cd_total_07 +
##   class_a:green_rating + empl_gr:class_a

## Warning: attempting model selection on an essentially perfect fit is
## nonsense

```

```

##              Df Sum of Sq    RSS      AIC
## - class_a:green_rating 1         0      0 -438769
## <none>                     0      0 -438768
## - class_a:util_index    1         0      0 -438766
## - empl_gr:class_a       1         0      0 -438764
## - LEED                  1         0      0 -438763
## - CS_PropertyID         1         0      0 -438755
## - class_a:cd_total_07   1         0      0 -438752
## - Electricity_Costs     1         0      0 -438740
## - leasing_rate          1         0      0 -438739
## - renovated             1         0      0 -438725
## - Energystar            1         0      0 -438718
## - hd_total07            1         0      0 -438717
## - stories               1         0      0 -438714
## - Precipitation         1         0      0 -438658
## - class_b               1         0      0 -438632
## - net                   1         0      0 -438584
## - amenities             1         0      0 -438523
## - Gas_Costs             1         0      0 -438521
## - age                   1         0      0 -438516
## - size                  1         0      0 -438466
## - cluster               1         0      0 -438171
## - cluster_rent          1    411193 411193   30436
## - Rent_Diff             1    680136 680136   34263
##
## Step:  AIC=-438769.7
## Rent ~ CS_PropertyID + cluster + size + empl_gr + leasing_rate +
##        stories + age + renovated + class_a + class_b + LEED + Energystar +
##        green_rating + net + amenities + cd_total_07 + hd_total07 +
##        Precipitation + Gas_Costs + Electricity_Costs + cluster_rent +
##        Rent_Diff + util_index + class_a:util_index + class_a:cd_total_07 +
##        empl_gr:class_a

## Warning: attempting model selection on an essentially perfect fit is
## nonsense

```

```

##              Df Sum of Sq    RSS      AIC
## <none>                     0 -438770
## - green_rating            1         0      0 -438768
## - class_a:util_index      1         0      0 -438768
## - empl_gr:class_a         1         0      0 -438766
## - LEED                    1         0      0 -438765
## - CS_PropertyID           1         0      0 -438757
## - class_a:cd_total_07     1         0      0 -438753
## - Electricity_Costs       1         0      0 -438742
## - leasing_rate            1         0      0 -438741
## - renovated               1         0      0 -438727
## - Energystar              1         0      0 -438720
## - hd_total07              1         0      0 -438719
## - stories                 1         0      0 -438715

```

```
## - Precipitation      1      0      0 -438660
## - class_b           1      0      0 -438634
## - net                1      0      0 -438591
## - amenities          1      0      0 -438525
## - Gas_Costs          1      0      0 -438523
## - age                1      0      0 -438518
## - size               1      0      0 -438469
## - cluster            1      0      0 -438173
## - cluster_rent       1    411770 411770   30443
## - Rent_Diff          1    680334 680334   34262
```

```
summary(lm.fit3)
```

```
##
## Call:
## lm(formula = Rent ~ CS_PropertyID + cluster + size + empl_gr +
##      leasing_rate + stories + age + renovated + class_a + class_b +
##      LEED + Energystar + green_rating + net + amenities + cd_total_07 +
##      hd_total07 + Precipitation + Gas_Costs + Electricity_Costs +
##      cluster_rent + Rent_Diff + util_index + class_a:util_index +
##      class_a:cd_total_07 + empl_gr:class_a, data = gbuild_sub)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.219e-12 -1.120e-14  1.300e-15  1.020e-14  2.559e-11
##
## Coefficients:
##              Estimate Std. Error  t value Pr(>|t|)
## (Intercept)   1.659e-13  7.400e-14  2.242e+00  0.02496 *
## CS_PropertyID  2.574e-20  5.408e-21  4.759e+00  1.98e-06 ***
## cluster       -2.849e-17  9.067e-18 -3.142e+00  0.00168 **
## size          -3.581e-20  2.088e-20 -1.715e+00  0.08636 .
## empl_gr       -3.208e-16  7.515e-16 -4.270e-01  0.66951
## leasing_rate   4.159e-16  2.150e-16  1.935e+00  0.05304 .
## stories        2.495e-16  5.125e-16  4.870e-01  0.62645
## age            8.009e-17  1.537e-16  5.210e-01  0.60240
## renovated      4.873e-16  8.301e-15  5.900e-02  0.95319
## class_a       -1.575e-14  2.017e-14 -7.810e-01  0.43503
## class_b       -6.861e-15  1.115e-14 -6.150e-01  0.53841
## LEED           2.681e-14  1.128e-13  2.380e-01  0.81219
## Energystar     6.321e-14  1.202e-13  5.260e-01  0.59899
## green_rating   -1.755e-14  1.209e-13 -1.450e-01  0.88457
## net            -2.121e-14  1.886e-14 -1.125e+00  0.26073
## amenities      7.999e-15  8.073e-15  9.910e-01  0.32183
## cd_total_07    2.087e-17  1.892e-17  1.103e+00  0.27014
## hd_total07     7.712e-18  7.788e-18  9.900e-01  0.32212
## Precipitation  -5.167e-16  5.193e-16 -9.950e-01  0.31977
## Gas_Costs      4.967e-12  4.336e-12  1.145e+00  0.25209
## Electricity_Costs -3.902e-13  1.156e-12 -3.370e-01  0.73576
## cluster_rent   1.000e+00  4.617e-16  2.166e+15 < 2e-16 ***
```

```
## Rent_Diff          1.000e+00  3.592e-16  2.784e+15  < 2e-16 ***
## util_index         -5.886e-16  5.952e-16 -9.890e-01  0.32279
## class_a:util_index -6.352e-16  2.957e-16 -2.148e+00  0.03176 *
## class_a:cd_total_07 5.244e-17  1.180e-17  4.443e+00  8.99e-06 ***
## empl_gr:class_a     -2.467e-15  9.796e-16 -2.519e+00  0.01180 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.962e-13 on 7579 degrees of freedom
## (73 observations deleted due to missingness)
## Multiple R-squared:  1, Adjusted R-squared:  1
## F-statistic: 7.666e+29 on 26 and 7579 DF, p-value: < 2.2e-16
```

```
confint(lm.fit3)
```

```
##              2.5 %          97.5 %
## (Intercept)  2.088064e-14  3.110011e-13
## CS_PropertyID 1.513610e-20  3.633968e-20
## cluster      -4.626273e-17 -1.071480e-17
## size         -7.674500e-20  5.118619e-21
## empl_gr      -1.793931e-15  1.152382e-15
## leasing_rate -5.447144e-18  8.373433e-16
## stories      -7.551991e-16  1.254132e-15
## age          -2.212730e-16  3.814592e-16
## renovated    -1.578538e-14  1.676003e-14
## class_a      -5.529029e-14  2.379531e-14
## class_b      -2.872016e-14  1.499860e-14
## LEED         -1.943457e-13  2.479583e-13
## Energystar   -1.724164e-13  2.988363e-13
## green_rating -2.545419e-13  2.194381e-13
## net          -5.818189e-14  1.575771e-14
## amenities    -7.827249e-15  2.382491e-14
## cd_total_07  -1.622360e-17  5.795607e-17
## hd_total07    -7.555322e-18  2.297868e-17
## Precipitation -1.534598e-15  5.012440e-16
## Gas_Costs     -3.533754e-12  1.346727e-11
## Electricity_Costs -2.656771e-12  1.876338e-12
## cluster_rent  1.000000e+00  1.000000e+00
## Rent_Diff     1.000000e+00  1.000000e+00
## util_index    -1.755330e-15  5.782299e-16
## class_a:util_index -1.214864e-15 -5.546056e-17
## class_a:cd_total_07 2.930674e-17 7.558264e-17
## empl_gr:class_a -4.387713e-15 -5.469624e-16
```

Surprisingly LEED remained in the model, however we fail to reject the null hypothesis that it is significant at the 95% level.

First we will begin by breaking the important continuous variables into manageable buckets. This will also serve us well to see the distribution of buildings across different ranges of values and setup our cross tab tables coming up. We will also limit the data to

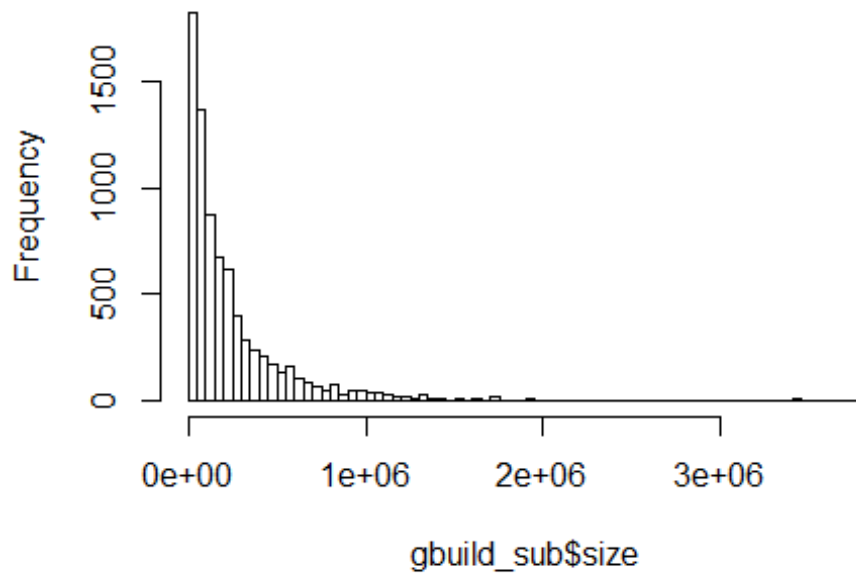
cities with positive employment growth, since Austin has one of the best economies in the country.

```
gbuild_sub$sizeCategory = cut(gbuild_sub$size, breaks = c(rep(0:20)*200000))
gbuild_sub$storiesCategory = cut(gbuild_sub$stories, breaks = c(rep(0:12)*10)
)
gbuild_sub$empl_grCategory = cut(gbuild_sub$empl_gr, breaks = c(rep(0:6)))
gbuild_sub$ageCategory = cut(gbuild_sub$age, breaks = c(rep(0:10)*20))
gbuild_sub$Electricity_CostsCategory = cut(gbuild_sub$Electricity_Costs,
                                           breaks = c(seq(0.00,0.07, by=0.01)
))
gbuild_sub$Gas_CostsCategory = cut(gbuild_sub$Gas_Costs,
                                   breaks = c(seq(0.00,0.03, by=0.005)
))
gbuild_sub$total_dd_07Category = cut(gbuild_sub$total_dd_07,
                                     breaks = c(seq(0.00,9000, by=2000)
))
gbuild_sub$cd_total_07Category = cut(gbuild_sub$cd_total_07,
                                     breaks = c(seq(0.00,6000, by=600))
)
attach(gbuild_sub)

## The following objects are masked from gbuild:
##
##      age, amenities, cd_total_07, class_a, class_b, cluster,
##      cluster_rent, CS_PropertyID, Electricity_Costs, empl_gr,
##      Energystar, Gas_Costs, green_rating, hd_total07, leasing_rate,
##      LEED, net, Precipitation, renovated, Rent, size, stories,
##      total_dd_07

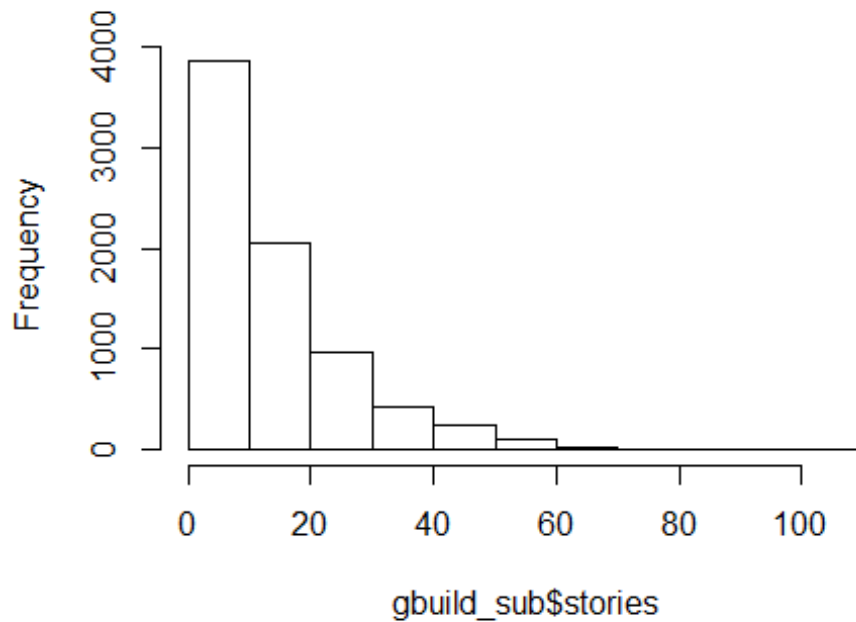
hist(gbuild_sub$size, breaks=75) #Good dispersion in target range
```


Histogram of gbuild_sub\$size



```
hist(gbuild_sub$stories) #Good dispersion in target range
```

Histogram of gbuild_sub\$stories




There appears to be a good distribution of buildings over the top predictors, so I'm not concerned about extrapolating outside of the observed ranges.

Next I inspect the theory that lower utility costs are the driver of higher rent prices in green buildings. In order to do this I create a feature called `util_index`, which is the sum of the products of gas costs and heating days and electric costs and cooling days. This feature will allow us to measure the expense of HVAC in a single variable.

We begin by examining rents for green and non-green buildings as two different series across the full range of `util_index`:

```
gbuild_sub$util_index = gbuild_sub$hd_total07*gbuild_sub$Gas_Costs +
  gbuild_sub$cd_total_07 * gbuild_sub$Electricity_Costs
#Bucket util_index
gbuild_sub$util_indexCategory = cut(gbuild_sub$util_index,
                                   breaks = c(seq(0.00,200, by=25)))
#Util index negatively correlated with rent; Higher utilities = Lower rent
plot(Rent~util_index,data = gbuild_sub, col="blue",pch=16)
```



```
#Normalize util_index in case we need it
gbuild_sub$util_index_norm = gbuild_sub$util_index/mean(gbuild_sub$util_index
)
```

```
summary(gbuild_sub)
```

##	CS_PropertyID	cluster	size	empl_gr
##	Min. : 1	Min. : 1.0	Min. : 2378	Min. : -24.950
##	1st Qu.: 157426	1st Qu.: 272.0	1st Qu.: 52000	1st Qu.: 1.740
##	Median : 313238	Median : 479.0	Median : 132417	Median : 1.970
##	Mean : 435335	Mean : 590.1	Mean : 239465	Mean : 3.188

```

## 3rd Qu.: 440780    3rd Qu.:1044.0    3rd Qu.: 302375    3rd Qu.:  2.380
## Max.    :6208103    Max.    :1230.0    Max.    :3781045    Max.    : 67.780
##
##                               NA's    :73
##      Rent      leasing_rate      stories      age
## Min.    : 2.98    Min.    : 10.68    Min.    :  1.00    Min.    :  0.00
## 1st Qu.: 19.50    1st Qu.: 79.51    1st Qu.:  4.00    1st Qu.: 23.00
## Median : 25.29    Median : 90.24    Median : 10.00    Median : 34.00
## Mean    : 28.59    Mean    : 84.88    Mean    : 13.83    Mean    : 47.04
## 3rd Qu.: 34.20    3rd Qu.: 96.66    3rd Qu.: 20.00    3rd Qu.: 79.00
## Max.    :250.00    Max.    :100.00    Max.    :110.00    Max.    :187.00
##
##      renovated      class_a      class_b      LEED
## Min.    :0.0000    Min.    :0.0000    Min.    :0.0000    Min.    :0.000000
## 1st Qu.:0.0000    1st Qu.:0.0000    1st Qu.:0.0000    1st Qu.:0.000000
## Median :0.0000    Median :0.0000    Median :0.0000    Median :0.000000
## Mean    :0.3814    Mean    :0.4083    Mean    :0.4587    Mean    :0.007032
## 3rd Qu.:1.0000    3rd Qu.:1.0000    3rd Qu.:1.0000    3rd Qu.:0.000000
## Max.    :1.0000    Max.    :1.0000    Max.    :1.0000    Max.    :1.000000
##
##      Energystar      green_rating      net      amenities
## Min.    :0.00000    Min.    :0.00000    Min.    :0.00000    Min.    :0.000
## 1st Qu.:0.00000    1st Qu.:0.00000    1st Qu.:0.00000    1st Qu.:0.000
## Median :0.00000    Median :0.00000    Median :0.00000    Median :1.000
## Mean    :0.08295    Mean    :0.08907    Mean    :0.03555    Mean    :0.538
## 3rd Qu.:0.00000    3rd Qu.:0.00000    3rd Qu.:0.00000    3rd Qu.:1.000
## Max.    :1.00000    Max.    :1.00000    Max.    :1.00000    Max.    :1.000
##
##      cd_total_07      hd_total07      total_dd_07      Precipitation
## Min.    : 39    Min.    :  0    Min.    :2103    Min.    :10.46
## 1st Qu.: 684    1st Qu.:1419    1st Qu.:2869    1st Qu.:22.71
## Median : 966    Median :2739    Median :4979    Median :23.16
## Mean    :1217    Mean    :3440    Mean    :4657    Mean    :31.10
## 3rd Qu.:1620    3rd Qu.:4796    3rd Qu.:6413    3rd Qu.:43.89
## Max.    :5240    Max.    :7200    Max.    :8244    Max.    :58.02
##
##      Gas_Costs      Electricity_Costs      cluster_rent      Rent_Diff
## Min.    :0.009487    Min.    :0.01780    Min.    : 9.00    Min.    : -45.9150
## 1st Qu.:0.010296    1st Qu.:0.02330    1st Qu.:20.25    1st Qu.: -2.9100
## Median :0.010296    Median :0.03274    Median :25.20    Median :  0.0000
## Mean    :0.011329    Mean    :0.03095    Mean    :27.60    Mean    :  0.9903
## 3rd Qu.:0.011816    3rd Qu.:0.03781    3rd Qu.:34.15    3rd Qu.:  3.3300
## Max.    :0.028914    Max.    :0.06280    Max.    :71.44    Max.    :191.2800
##
##      util_index      sizeCategory      storiesCategory      empl_grCategory
## Min.    : 33.12    (0,2e+05]      :4746    (0,10] :3867    (0,1]: 916
## 1st Qu.: 40.47    (2e+05,4e+05] :1530    (10,20]:2049    (1,2]:2970
## Median : 78.53    (4e+05,6e+05] : 672    (20,30]: 963    (2,3]:2307
## Mean    : 75.78    (6e+05,8e+05] : 297    (30,40]: 420    (3,4]: 624
## 3rd Qu.: 96.17    (8e+05,1e+06] : 192    (40,50]: 245    (4,5]: 434
## Max.    :304.83    (1e+06,1.2e+06]: 117    (50,60]: 106    (5,6]: 98


```

```

##      (Other)      : 125  (Other): 29  NA's : 330
##      ageCategory Electricity_CostsCategory Gas_CostsCategory
## (20,40] :3183 (0,0.01] : 0 (0,0.005] : 0
## (0,20] :1287 (0.01,0.02]: 826 (0.005,0.01]: 554
## (80,100]:1250 (0.02,0.03]:2917 (0.01,0.015]:6993
## (40,60] : 886 (0.03,0.04]:3535 (0.015,0.02]: 45
## (60,80] : 553 (0.04,0.05]: 314 (0.02,0.025]: 0
## (Other) : 503 (0.05,0.06]: 0 (0.025,0.03]: 87
## NA's : 17 (0.06,0.07]: 87
##      total_dd_07Category cd_total_07Category util_indexCategory
## (0,2e+03] : 0 (600,1.2e+03] :3523 (25,50] :2872
## (2e+03,4e+03]:3044 (0,600] :1667 (75,100] :2428
## (4e+03,6e+03]:2087 (1.8e+03,2.4e+03]: 823 (100,125]:1406
## (6e+03,8e+03]:2308 (1.2e+03,1.8e+03]: 695 (50,75] : 626
## NA's : 240 (2.4e+03,3e+03] : 424 (125,150]: 217
## (4.8e+03,5.4e+03]: 259 (Other) : 43
## (Other) : 288 NA's : 87
##      util_index_norm
##      Min. :0.4370
##      1st Qu.:0.5341
##      Median :1.0364
##      Mean :1.0000
##      3rd Qu.:1.2692
##      Max. :4.0228
##
g_green = subset(gbuild_sub, green_rating==1)
g_green = g_green[complete.cases(g_green),]
g_ngreen = subset(gbuild_sub, green_rating==0)
g_ngreen = g_ngreen[complete.cases(g_green),]


plot(Rent~util_index,data = g_ngreen, col="blue",pch=16)
points(Rent~util_index,data = g_green, col="red",pch=4)

```




Well there isn't much to take from this graph. Let's try to do the same thing across total-dd_07-days:

```
plot(Rent~total_dd_07,data = g_ngreen, col="blue",pch=16)
points(Rent~total_dd_07,data = g_green, col="red",pch=4)
```



```
plot(Rent~cd_total_07,data = g_ngreen, col="blue",pch=16)
points(Rent~cd_total_07,data = g_green, col="red",pch=4)
```



It appears green buildings are highly concentrated in milder climates. Lets look at a cross tab of green building frequencies by util_index and class_a.

```
freq =xtabs(~green_rating+class_a+util_indexCategory, data = gbuild_sub)
freq

## , , util_indexCategory = (0,25]
##
##           class_a
## green_rating    0    1
##           0      0    0
##           1      0    0
##
## , , util_indexCategory = (25,50]
##
##           class_a
## green_rating    0    1
##           0 1624  951
##           1   70  227
##
## , , util_indexCategory = (50,75]
##
##           class_a
## green_rating    0    1
##           0  400  191
```

```

##           1    7   28
##
## , , util_indexCategory = (75,100]
##
##           class_a
## green_rating    0    1
##           0 1518  704
##           1   42  164
##
## , , util_indexCategory = (100,125]
##
##           class_a
## green_rating    0    1
##           0  714  588
##           1   12   92
##
## , , util_indexCategory = (125,150]
##
##           class_a
## green_rating    0    1
##           0   94   94
##           1    1   28
##
## , , util_indexCategory = (150,175]
##
##           class_a
## green_rating    0    1
##           0    0    0
##           1    0    0
##
## , , util_indexCategory = (175,200]
##
##           class_a
## green_rating    0    1
##           0   28    6
##           1    4    5

```

The table above seems to indicate green buildings are highly concentrated in class a buildings. That would be a good reason why they appear to rent for more money. The table shows 3 splits of the util_index: 0-75,75-150,150-225

For those respective bins, 255 of the 332 green buildings or 76.8%, are found in class a buildings. 284/339 or 83.8% of green buildings are class a in the second bin. And 5/9 of the green buildings in the last bin are class A. It seems we've found something here. What if the higher rent prices for green buildings were a reflection of the class of the building instead of the green rating?

Lets dig deeper, and look at a few boxplots of rent by green rating for only class A buildings & non-class A buildings

```


gbuild_sub_A = subset(gbuild_sub, class_a == 1)
gbuild_sub_NotA = subset(gbuild_sub, class_a == 0)
boxplot(Rent ~ green_rating+util_indexCategory, data = gbuild_sub_A,
        xlab= "Green Rating followed by Util_Index",
        ylab="Rent",notch=TRUE,ylim=c(5,80),cex.axis=.7)

## Warning in bxp(structure(list(stats = structure(c(NA, NA, NA, NA, NA, NA,
:
## some notches went outside hinges ('box'): maybe set notch=FALSE

        title("Class A Rent Prices by (Green Rating and Util Index)")

```

Class A Rent Prices by (Green Rating and Util Inde



```

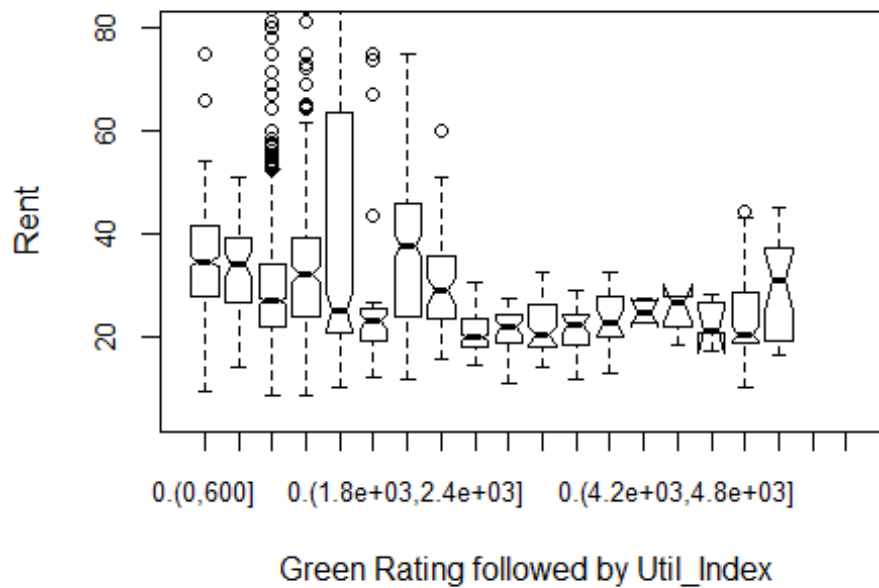
boxplot(Rent ~ green_rating+cd_total_07Category, data = gbuild_sub_A,
        xlab= "Green Rating followed by Util_Index",
        ylab="Rent",notch=TRUE,ylim=c(5,80),cex.axis=.8)

## Warning in bxp(structure(list(stats = structure(c(9.6, 28.2, 34.84, 41.5,
:
## some notches went outside hinges ('box'): maybe set notch=FALSE

        title("Class A Rent Prices by (Green Rating and Util Index)")


```


Class A Rent Prices by (Green Rating and Util Inde



```
#boxplot(Rent ~ green_rating+total_dd_07Category, data = gbuild_sub_A, xlab=
"Green Rating followed by total degree days", ylab="Rent",notch=TRUE,ylim=c(5
,80),cex.axis=.8)
boxplot(Rent ~ green_rating+util_indexCategory,
        data = gbuild_sub_NotA, xlab= "Green Rating followed by Util_Index",y
lab="Rent",
        notch=FALSE,ylim=c(5,80),cex.axis=.8)
title("Non-Class A Rent Prices by (Green Rating and Util Index)")
```

Non-Class A Rent Prices by (Green Rating and Util In



```
g_control1 = subset(gbuild_sub_A, net==1 & leasing_rate <= 80 & empl_gr > 0)
freq =xtabs(~green_rating+empl_grCategory+util_indexCategory, data = g_control1)
freq
```

```
## , , util_indexCategory = (0,25]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0      0      0      0      0      0
##           1      0      0      0      0      0
##
## , , util_indexCategory = (25,50]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0      0      3      3      0      0
##           1      0      2      0      0      0
##
## , , util_indexCategory = (50,75]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0      0      0      0      0      0
##           1      0      0      0      0      0
##
## , , util_indexCategory = (75,100]
```

```

##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0      0      6      0      0      0      2
##           1      0      0      1      1      0      0
##
## , , util_indexCategory = (100,125]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0      1      0      0      6      0      0
##           1      0      0      0      2      0      0
##
## , , util_indexCategory = (125,150]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0      0      0      7      0      0      0
##           1      0      0      1      0      0      0
##
## , , util_indexCategory = (150,175]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0      0      0      0      0      0      0
##           1      0      0      0      0      0      0
##
## , , util_indexCategory = (175,200]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0      0      0      0      0      0      0
##           1      0      0      0      0      0      0
##
rent_sum =xtabs(Rent~green_rating+empl_grCategory+util_indexCategory, data =
g_control1)
avg_rent = rent_sum/freq
avg_rent

## , , util_indexCategory = (0,25]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0
##           1
##
## , , util_indexCategory = (25,50]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]

```

```

##          0          28.33333 33.51333
##          1          30.31500
##
## , , util_indexCategory = (50,75]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##          0
##          1
##
## , , util_indexCategory = (75,100]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##          0          23.50000          25.00000
##          1          14.34000 19.25000
##
## , , util_indexCategory = (100,125]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##          0 23.65000          18.77167
##          1          20.92000
##
## , , util_indexCategory = (125,150]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##          0          17.84429
##          1          18.61000
##
## , , util_indexCategory = (150,175]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##          0
##          1
##
## , , util_indexCategory = (175,200]
##
##          empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##          0
##          1

```


At this point we can see green buildings are highly correlated and that some green buildings in certain utility_index bins do rent for a premium. The last table is particularly interesting. Here we can see that green buildings generally only sell for a premium in modest to high growth cities.

Net pricing wasn't indicated as a strong predictor but let's do the same exercise controlling for non-net leases and some other features correlated with price.

#Control for important lurking variables

```
g_control1 = subset(gbuild_sub_A, net==1 & leasing_rate <= 80 & empl_gr > 0)
g_control2 = subset(gbuild_sub_A, net==1 & leasing_rate <= 80 & empl_gr <= 0)
```

```
plot(gbuild$empl_gr)
```




```
boxplot(Rent ~ green_rating+util_indexCategory, data = g_control1,
        xlab= "Green Rating followed by Util_Index",
        ylab="Rent",notch=TRUE,ylim=c(5,80),cex.axis=.8)
```

```
## Warning in bxp(structure(list(stats = structure(c(NA, NA, NA, NA, NA, NA,
:
## some notches went outside hinges ('box'): maybe set notch=FALSE
```


```
title("Class A Rent Prices by (Green Rating and Util Index)")
```

Class A Rent Prices by (Green Rating and Util Inde



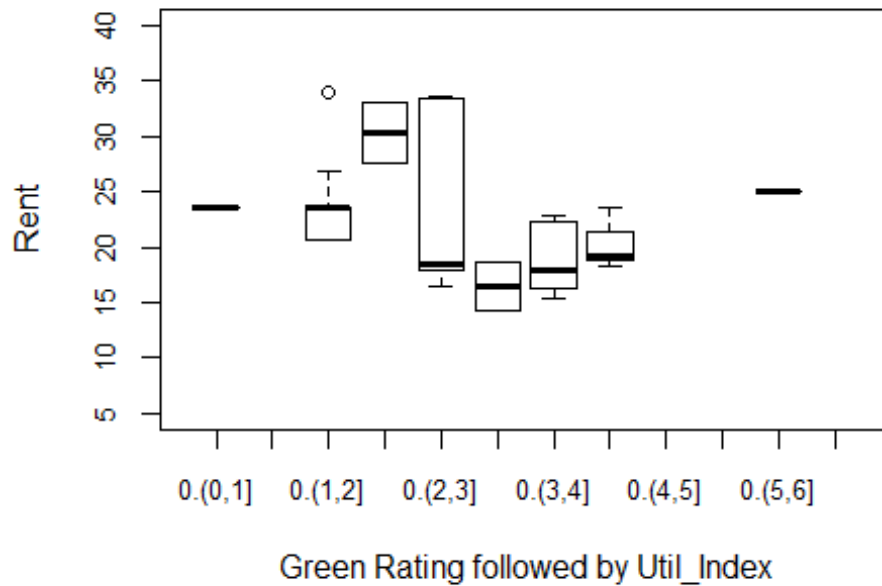
```
boxplot(Rent ~ green_rating+cd_total_07Category, data = g_control1,  
        xlab= "Green Rating followed by Util_Index",  
        ylab="Rent",ylim=c(5,40),cex.axis=.8)  
title("Class A Rent Prices by (Green Rating and Util Index)")
```

Class A Rent Prices by (Green Rating and Util Inde




```
boxplot(Rent ~ green_rating+empl_grCategory, data = g_control1,  
        xlab= "Green Rating followed by Util_Index",  
        ylab="Rent",ylim=c(5,40),cex.axis=.8)  
title("Class A Rent Prices by (Green Rating and Util Index)")
```

Class A Rent Prices by (Green Rating and Util Inde



```
plot(Rent_Diff~util_index,data = g_ngreen, col="blue",pch=16)  
points(Rent_Diff~util_index,data = g_green, col="red",pch=4)
```




```

freq =xtabs(~green_rating+class_a+util_indexCategory, data = gbuild_sub)
rent_sum =xtabs(Rent~green_rating+class_a+util_indexCategory, data = gbuild_sub)
avg_rent = rent_sum/freq
avg_rent

## , , util_indexCategory = (0,25]
##
##           class_a
## green_rating 0 1
##           0
##           1
##
## , , util_indexCategory = (25,50]
##
##           class_a
## green_rating 0 1
##           0 30.76389 36.94904
##           1 31.60100 37.39678
##
## , , util_indexCategory = (50,75]
##
##           class_a
## green_rating 0 1
##           0 23.24820 31.24958
##           1 24.31571 31.23036
##
## , , util_indexCategory = (75,100]
##
##           class_a
## green_rating 0 1
##           0 19.39635 21.75456
##           1 19.52500 22.79463
##
## , , util_indexCategory = (100,125]
##
##           class_a
## green_rating 0 1
##           0 31.18513 40.77844
##           1 22.53083 30.21543
##
## , , util_indexCategory = (125,150]
##
##           class_a
## green_rating 0 1
##           0 26.40766 31.10670
##           1 20.90000 31.30321
##
## , , util_indexCategory = (150,175]
##

```

```

##           class_a
## green_rating 0 1
##           0
##           1
##
## , , util_indexCategory = (175,200]
##
##           class_a
## green_rating      0      1
##           0 20.44107 25.05500
##           1 24.12000 23.05200

gbuild_notnet = subset(gbuild_sub,net==0)
gbuild_net = subset(gbuild_sub,net==1)

g_cntrl_notnet = subset(gbuild_notnet, class_a == 1 | class_b == 1 & age <= 3
0 &
                        (empl_gr >= 1 & empl_gr <= 3) &
                        (stories >= 5 & stories <= 25) &
                        (size <= 300000 & size >= 200000))

freq =xtabs(~green_rating+empl_grCategory+class_a, data = gbuild_sub)
freq

## , , class_a = 0
##
##           empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0   491  1925  1166   331   261    52
##           1    20    39    46     9    13     1
##
## , , class_a = 1
##
##           empl_grCategory
## green_rating (0,1] (1,2] (2,3] (3,4] (4,5] (5,6]
##           0   359   876   877   214   119    35
##           1    46   130   218    70    41    10

rent_sum =xtabs(Rent~green_rating+empl_grCategory+class_a, data = gbuild_sub)
avg_rent = rent_sum/freq
avg_rent

## , , class_a = 0
##
##           empl_grCategory
## green_rating  (0,1]  (1,2]  (2,3]  (3,4]  (4,5]  (5,6]
##           0 31.16255 27.85325 25.59782 19.44486 17.98046 18.21058
##           1 26.38250 29.19846 29.17696 19.60889 18.52154 10.51000
##
## , , class_a = 1
##

```

```

##                empl_grCategory
## green_rating  (0,1]  (1,2]  (2,3]  (3,4]  (4,5]  (5,6]
##                0 35.87744 36.09933 32.41568 23.99907 20.45521 21.53857
##                1 32.96804 31.20923 35.11009 24.44957 21.04268 22.44500

freq =xtabs(~green_rating+total_dd_07Category, data = gbuild_sub)
rent_sum =xtabs(Rent~green_rating+total_dd_07Category, data = gbuild_sub)
avg_rent = rent_sum/freq
avg_rent

##                total_dd_07Category
## green_rating (0,2e+03] (2e+03,4e+03] (4e+03,6e+03] (6e+03,8e+03]
##                0      32.41796      27.41713      26.10320
##                1      35.60505      25.70507      25.92176

freq =xtabs(~green_rating+net+util_indexCategory, data = gbuild_sub)
rent_sum =xtabs(Rent~green_rating+net+util_indexCategory, data = gbuild_sub)
avg_rent = rent_sum/freq
avg_rent

## , , util_indexCategory = (0,25]
##
##                net
## green_rating 0 1
##                0
##                1
##
## , , util_indexCategory = (25,50]
##
##                net
## green_rating 0      1
##                0 33.18523 26.26608
##                1 36.10089 31.93600
##
## , , util_indexCategory = (50,75]
##
##                net
## green_rating 0      1
##                0 26.02386 17.39692
##                1 30.32531 24.75000
##
## , , util_indexCategory = (75,100]
##
##                net
## green_rating 0      1
##                0 20.16530 19.48324
##                1 22.21375 20.95214
##
## , , util_indexCategory = (100,125]
##
##                net

```

```

## green_rating      0      1
##                0 35.73954 31.22422
##                1 29.80478 26.26857
##
## , , util_indexCategory = (125,150]
##
##                net
## green_rating      0      1
##                0 29.23134 19.32667
##                1 31.82778 19.02000
##
## , , util_indexCategory = (150,175]
##
##                net
## green_rating 0 1
##                0
##                1
##
## , , util_indexCategory = (175,200]
##
##                net
## green_rating      0      1
##                0 21.14091 25.03000
##                1 23.52667

freq

## , , util_indexCategory = (0,25]
##
##                net
## green_rating      0      1
##                0      0      0
##                1      0      0
##
## , , util_indexCategory = (25,50]
##
##                net
## green_rating      0      1
##                0 2524      51
##                1 292      5
##
## , , util_indexCategory = (50,75]
##
##                net
## green_rating      0      1
##                0 578      13
##                1 32      3
##
## , , util_indexCategory = (75,100]
##

```

```
##           net
## green_rating  0    1
##           0 2151   71
##           1  192   14
##
## , , util_indexCategory = (100,125]
##
##           net
## green_rating  0    1
##           0 1238   64
##           1   90   14
##
## , , util_indexCategory = (125,150]
##
##           net
## green_rating  0    1
##           0  179    9
##           1   27    2
##
## , , util_indexCategory = (150,175]
##
##           net
## green_rating  0    1
##           0    0    0
##           1    0    0
##
## , , util_indexCategory = (175,200]
##
##           net
## green_rating  0    1
##           0   33    1
##           1    9    0
```

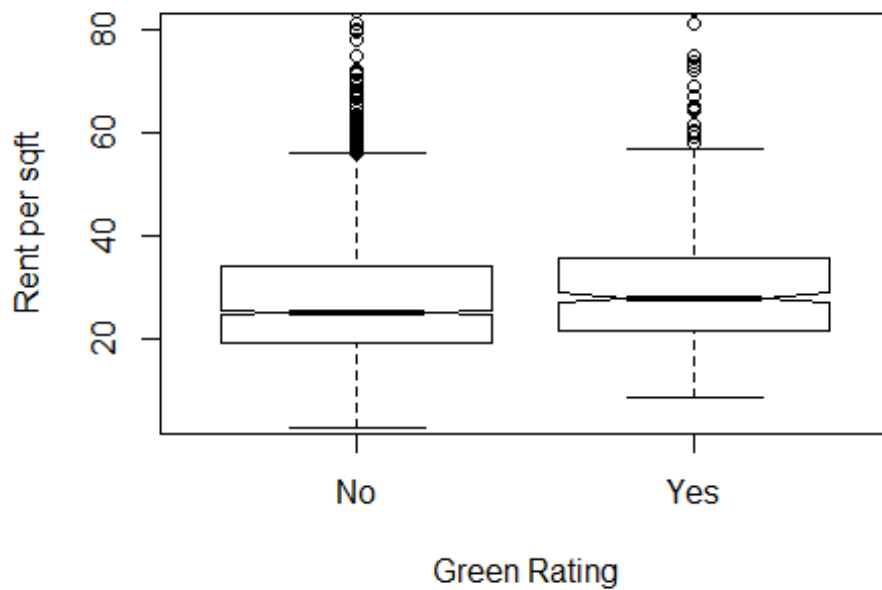
```
plot(cluster_rent~util_index, data = subset(gbuild_sub,green_rating==1),col="blue")
points(cluster_rent~util_index, data = subset(gbuild_sub,green_rating==0),col="red")
```

```
plot(Rent_norm~util_index, data = subset(gbuild_sub,green_rating==1),col="blue",pch=16)
points(Rent_norm~util_index, data = subset(gbuild_sub,green_rating==0),col="red",
pch=4)
```


```
...
```

Now lets examine rent by utility costs, after we control for some features. Let's assume the building will be class A, with median employment growth (2), roughly 250,000 sqft, and less than 10 years old.

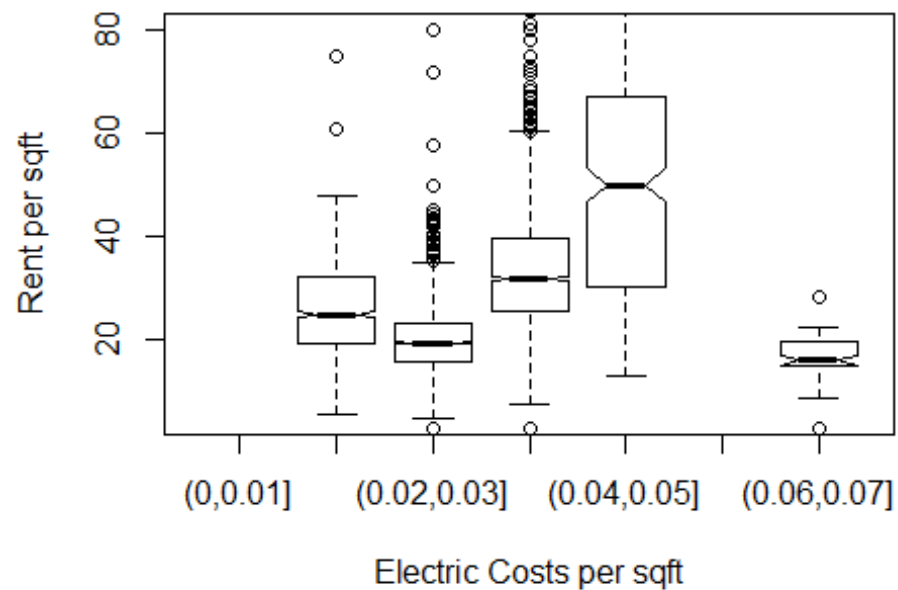
```
boxplot(Rent ~ green_rating, data = gbuild_notnet, names= c("No", "Yes"),title
="Non-Net Leases", xlab= "Green Rating", ylab="Rent per sqft",notch=TRUE, yli
m=c(5,80))
```



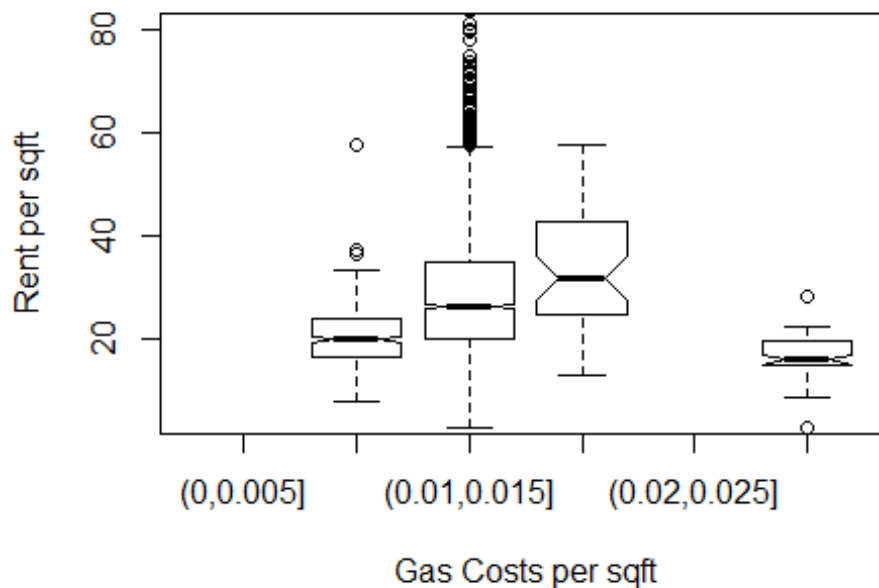
```
boxplot(Rent ~ green_rating, data = gbuild_net, names= c("No","Yes"),title="Net Leases", xlab= "Green Rating", ylab="Rent per sqft",notch=TRUE, ylim=c(5,80))
```



```
boxplot(Rent ~ Electricity_CostsCategory, data = gbuild_notnet, xlab= "Electric Costs per sqft", ylab="Rent per sqft",notch=TRUE,ylim=c(5,80))
```



```
boxplot(Rent ~ Gas_CostsCategory, data = gbuild_notnet, xlab= "Gas Costs per sqft", ylab="Rent per sqft",notch=TRUE,ylim=c(5,80))
```



```
xtabs(~green_rating+net,data = gbuild_sub)
```

```
##           net
## green_rating  0    1
##           0 6761  234
##           1  645   39
```

Conclusion -

In conclusion, we have shown how risky and unreliable the former analysts recommendations were. By not using regression to control for the other features the recommendations was wreckless. We have attempted to isolate the effect of the green buildings outside of regression. While we weren't able to isolate the effect of green buildings completely, we believe its quite evident its highly correlated with other features that drive up rent, such as class a buildings and high employment growth cities. At a minimum we have shown that green buildings effect is not consistent throughout the data set, and that its unwise to generalize.

Bootstrapping :

The value at risk and returns of each portfolio gives us a measure of how "safe" or "risky" an asset is.

```
suppressMessages(library(mosaic))
suppressMessages(library(fImport))
suppressMessages(library(foreach))
```



```

mystocks = c("SPY","TLT","LQD","EEM","VNO")
myprices = yahooSeries(mystocks, from='2010-01-01', to='2016-07-30')

# A helper function for calculating percent returns from a Yahoo Series
YahooPricesToReturns = function(series) {
  mycols = grep('Adj.Close', colnames(series))
  closingprice = series[,mycols]
  N = nrow(closingprice)
  percentreturn = as.data.frame(closingprice[2:N,]) / as.data.frame(closing
price[1:(N-1),]) - 1
  mynames = strsplit(colnames(percentreturn), '.', fixed=TRUE)
  mynames = lapply(mynames, function(x) return(paste0(x[1], ".PctReturn")))
  colnames(percentreturn) = mynames
  as.matrix(na.omit(percentreturn))
}

myreturns = YahooPricesToReturns(myprices)

```

marshals appropriate evidence to characterize the risk/return properties of the five major asset classes listed above.

We will use bootstrap sampling to calculate the returns for each asset. The code below does bootstrapping for SPY alone. Similarly we implement the code for all assets.

```

sim_SPY = foreach(i=1:5000, .combine='rbind') %do% {
  totalwealth = 100000
  n_days = 20
  weights_even = c(1.0, 0.0, 0.0, 0.0, 0.0)
  holdings = weights_even * totalwealth
  wealthtracker = rep(0, n_days)
  for(today in 1:n_days) {
    return.today = resample(myreturns, 1, orig.ids=FALSE)
    holdings = holdings + holdings*return.today
    totalwealth = sum(holdings)
    wealthtracker[today] = totalwealth
    holdings = weights_even * totalwealth
  }
  wealthtracker
}

```

The average returns for SPY over 20 days is

```
mean(sim_SPY[,n_days])
```

```
## [1] 101076.9
```

5% value at risk for SPY is :

```
quantile(sim_SPY[,n_days], 0.05) - 100000
```

```
##          5%
## -6183.525
```

outlines your choice of the "safe" and "aggressive" portfolios.

We derived a table like the one below to identify the assets as safe and aggressive based on their loss at risk and average returns.

uses bootstrap resampling to estimate the 4-week (20 trading day) value at risk of each of your three portfolios at the 5% level

Even split portfolio :

```
sim_even = foreach(i=1:500, .combine='rbind') %do% {
  totalwealth = 100000
  n_days = 20
  weights_even = c(0.2, 0.2, 0.2, 0.2, 0.2)
  holdings = weights_even * totalwealth
  wealthtracker = rep(0, n_days)
  for(today in 1:n_days) {
    return.today = resample(myreturns, 1, orig.ids=FALSE)
    holdings = holdings + holdings*return.today
    totalwealth = sum(holdings)
    wealthtracker[today] = totalwealth
    holdings = weights_even * totalwealth
  }
  wealthtracker
}
```

Average return for even split portfolio is

```
return_even <- mean(sim_even[,n_days])
```

5% value at risk for even split portfolio

```
risk_even <- quantile(sim_even[,n_days], 0.05) - 100000
```

Safe portfolio - The safe portfolio will use the safest assets - SPY, TLT and LQD (at least 3 classes required)The safe assets are those that have low risk. We are choosing to invest about 80% of our wealth into SPY because SPY has the highest returns among the three and has medium to low risk :

```
sim_safe = foreach(i=1:500, .combine='rbind') %do% {
  totalwealth = 100000
  n_days = 20
  weights_even = c(0.8, 0.1, 0.1, 0.0, 0.0)
  holdings = weights_even * totalwealth
  wealthtracker = rep(0, n_days)
  for(today in 1:n_days) {
    return.today = resample(myreturns, 1, orig.ids=FALSE)
```

```

    holdings = holdings + holdings*return.today
    totalwealth = sum(holdings)
    wealthtracker[today] = totalwealth
    holdings = weights_even * totalwealth
  }
  wealthtracker
}

```

Average return for safe split portfolio is :

```
return_safe <- mean(sim_safe[,n_days])
```

5% value at risk for safe split portfolio :

```
risk_safe <- quantile(sim_safe[,n_days], 0.05) - 100000
```

Aggressive portfolio : In our 'Aggressive portfolio', we have chosen the assets that give the highest returns irrespective of the risk involved. EEM, VNQ are the two assets that gave us the highest returns. So, our aggressive portfolio includes EEM and VNQ. We are choosing to invest in EEM and VNQ in the ratio 3:7 because VNQ offers higher returns than EEM and we want to maximize our returns.

```

sim_high = foreach(i=1:500, .combine='rbind') %do% {
  totalwealth = 100000
  n_days = 20
  weights_even = c(0.0, 0.0, 0.0, 0.3, 0.7)
  holdings = weights_even * totalwealth
  wealthtracker = rep(0, n_days)
  for(today in 1:n_days) {
    return.today = resample(myreturns, 1, orig.ids=FALSE)
    holdings = holdings + holdings*return.today
    totalwealth = sum(holdings)
    wealthtracker[today] = totalwealth
    holdings = weights_even * totalwealth
  }
  wealthtracker
}

```

Average return for aggressive portfolio is

```
return_aggressive <- mean(sim_high[,n_days])
```

5% value at risk for aggressive portfolio :

```
risk_aggressive <- quantile(sim_high[,n_days], 0.05) - 100000
```

compares the results for each portfolio in a way that would allow the reader to make an intelligent decision among the three options.

Conclusion

Average returns over a 20 day period for the three portfolios :

Even :

```
return_even
## [1] 100769.3
```

Safe :

```
return_safe
## [1] 101402
```

Aggressive :

```
return_aggressive
## [1] 101438.9
```

Loss at risk for the three portfolios :

Even :

```
risk_even
##          5%
## -4040.989
```

Safe :

```
risk_safe
##          5%
## -3946.606
```

Aggressive :

```
risk_aggressive
##          5%
## -7142.911
```

So, from the above estimations of risk and returns, if an investor is willing to be aggressive, then he stands to gain a lot in the returns and his loss at risk is also the highest among the three portfolios.

The safe portfolio does not yield higher returns than even portfolio and the loss at risk is also higher for safe portfolio as compared to the loss at risk value for even portfolio.

So, it is more beneficial to invest in an even portfolio.

Problem 3 :

Market segmentation

Initial Set-up and Loading the Data:

```
# Change to required path
```

```
library(flexclust)
```

```
## Loading required package: grid
```

```
## Loading required package: modeltools
```

```
## Loading required package: stats4
```

```
##
```

```
## Attaching package: 'modeltools'
```

```
## The following object is masked from 'package:RCurl':
```

```
##
```

```
##      clone
```

```
library(ggplot2)
```

```
library(reshape2)
```

```
library(corrplot)
```

```
library(corrgram)
```

```
mkt_seg = read.csv("C:/MSBA/James Scott Statistics/STA380-master/STA380-master/data/social_marketing.csv",header=T)
```

```
str(mkt_seg)
```

```
## 'data.frame':    7882 obs. of  37 variables:
```

```
## $ X                : Factor w/ 7882 levels "123pxkyqj","12grikctu",...: 372 0 2540 4096 596 3197 3609 4749 6518 7418 4917 ...
```

```
## $ chatter          : int  2 3 6 1 5 6 1 5 6 5 ...
```

```
## $ current_events   : int  0 3 3 5 2 4 2 3 2 2 ...
```

```
## $ travel           : int  2 2 4 2 0 2 7 3 0 4 ...
```

```
## $ photo_sharing    : int  2 1 3 2 6 7 1 6 1 4 ...
```

```
## $ uncategorized    : int  2 1 1 0 1 0 0 1 0 0 ...
```

```
## $ tv_film          : int  1 1 5 1 0 1 1 1 0 5 ...
```

```
## $ sports_fandom    : int  1 4 0 0 0 1 1 1 0 9 ...
```

```
## $ politics         : int  0 1 2 1 2 0 11 0 0 1 ...
```

```
## $ food             : int  4 2 1 0 0 2 1 0 2 5 ...
```

```
## $ family           : int  1 2 1 1 1 1 0 0 2 4 ...
```

```
## $ home_and_garden  : int  2 1 1 0 0 1 0 0 1 0 ...
```

```
## $ music            : int  0 0 1 0 0 1 0 2 1 1 ...
```

```
## $ news             : int  0 0 1 0 0 0 1 0 0 0 ...
```

```
## $ online_gaming    : int  0 0 0 0 3 0 0 1 2 1 ...
```

```
## $ shopping         : int  1 0 2 0 2 5 1 3 0 0 ...
```

```
## $ health_nutrition: int  17 0 0 0 0 0 1 1 22 7 ...
```

```
## $ college_uni      : int  0 0 0 1 4 0 1 0 1 4 ...
## $ sports_playing   : int  2 1 0 0 0 0 1 0 0 1 ...
## $ cooking          : int  5 0 2 0 1 0 1 10 5 4 ...
## $ eco              : int  1 0 1 0 0 0 0 0 2 1 ...
## $ computers        : int  1 0 0 0 1 1 1 1 1 2 ...
## $ business         : int  0 1 0 1 0 1 3 0 1 0 ...
## $ outdoors         : int  2 0 0 0 1 0 1 0 3 0 ...
## $ crafts           : int  1 2 2 3 0 0 0 1 0 0 ...
## $ automotive       : int  0 0 0 0 0 1 0 1 0 4 ...
## $ art              : int  0 0 8 2 0 0 1 0 1 0 ...
## $ religion          : int  1 0 0 0 0 0 1 0 0 13 ...
## $ beauty           : int  0 0 1 1 0 0 0 5 5 1 ...
## $ parenting        : int  1 0 0 0 0 0 0 1 0 3 ...
## $ dating           : int  1 1 1 0 0 0 0 0 0 0 ...
## $ school           : int  0 4 0 0 0 0 0 0 1 3 ...
## $ personal_fitness : int  11 0 0 0 0 0 0 0 12 2 ...
## $ fashion          : int  0 0 1 0 0 0 0 4 3 1 ...
## $ small_business   : int  0 0 0 0 1 0 0 0 1 0 ...
## $ spam             : int  0 0 0 0 0 0 0 0 0 0 ...
## $ adult            : int  0 0 0 0 0 0 0 0 0 0 ...
```


From looking at the various columns in the dataset, we decided to drop the columns spam and adult since they do not give us real insights into user preferences. In addition, we also combined the columns chatter and uncategorized into one since they represent the tweets that dont fit into any category.

```
mkt_seg_junk = mkt_seg[, -c(36,37)]
mkt_seg_junk$chatter = mkt_seg_junk$uncategorized + mkt_seg_junk$chatter
mkt_seg_junk = mkt_seg_junk[, -6] # Removing uncategorized

# Without the id column
mkt_seg_no_id = mkt_seg_junk[, -1]
```

To see if any of the variables are related, we plotted correlations. corrplot was used since it allows for easier and cleaner visualization of relationships.

```
# Looking at correlations between variables
corr_matrix = cor(mkt_seg_no_id)
corrplot(corr_matrix, type="lower", order="hclust")
```




From the corrplot, it seems like there are likely to be about 4-8 clusters.

To find the optimal number of clusters, we implemented the Elbow method for k means clustering after scaling and centering the data.


```
library(factoextra)
library(cluster)
library(NbClust)

# scaling before clustering
mkt_seg_scale <- scale(mkt_seg_no_id, center=TRUE, scale=TRUE)

set.seed(5)
# Calculating wss till k=15
k.max <- 15
data <- mkt_seg_scale
wss <- sapply(1:k.max,
              function(k){kmeans(data, k, nstart=10)$tot.withinss})
plot(1:k.max, wss,
     type="b", pch = 19, frame = FALSE,
     xlab="Number of clusters K",
     ylab="Total within-clusters sum of squares")
abline(v = 8, lty = 2)
```



```
# Cross checking with factoextra package  
fviz_nbclust(data, kmeans, method = "wss")
```




From the plots, the optimal number of clusters is 8. We chose 8 since it minimises wss to an acceptable value and will not have too many clusters that will it hard to interpret. We will use k means clustering with 8 clusters.

```
# K-means clustering
set.seed(10)
km_seg <- kmeans(mkt_seg_scale, 8, nstart = 30)

# k-means group number of each observation
clust_obs <- km_seg$cluster
table(clust_obs)

## clust_obs
##      1      2      3      4      5      6      7      8
## 794  440 3480  371  364  698  504 1231

# Visualize k-means clusters
fviz_cluster(km_seg, data = mkt_seg_scale, geom = "point",
              stand = FALSE, frame.type = "norm")
```



```
# Identifying where the centers of the clusters are
clusters_cent = km_seg$centers
imp_fact = t(clusters_cent)

# Separating by cluster and only taking important features
cluster_1 = imp_fact[which(abs(imp_fact[,1])>=0.4),1]
cluster_2 = imp_fact[which(abs(imp_fact[,2])>=0.4),2]
```

```

cluster_3 = imp_fact[which(abs(imp_fact[,3])>=0.4),3]
names(cluster_3) = c("photo_sharing") # since only 1 variable
cluster_4 = imp_fact[which(abs(imp_fact[,4])>=0.4),4]
cluster_5 = imp_fact[which(abs(imp_fact[,5])>=0.4),5]
cluster_6 = imp_fact[which(abs(imp_fact[,6])>=0.4),6]
cluster_7 = imp_fact[which(abs(imp_fact[,7])>=0.4),7]
cluster_8 = imp_fact[which(abs(imp_fact[,8])>=0.4),8]

```

Seeing how the clusters turned out
cluster_1

```

##          food health_nutrition          cooking          eco
##      0.4577940      2.2002089      0.4025928      0.5419657
##      outdoors personal_fitness
##      1.7114549      2.1673465

```

cluster_2

```

## sports_fandom      politics          news      automotive
##      0.6580063      1.2176757      2.6381607      2.5800077

```

cluster_3

```

## photo_sharing
##      -0.4056307

```

cluster_4

```

## online_gaming      college_uni sports_playing
##      3.497844      3.267583      2.149292

```

cluster_5

```

##      travel      politics          news      computers      business
##      3.2262928      3.0806307      1.1301224      2.8876695      0.5512909
## small_business
##      0.4158731

```

cluster_6

```

## sports_fandom      food          family      crafts      religion
##      2.0833636      1.8417506      1.5028191      0.7294984      2.2793038
##      parenting      school
##      2.1526532      1.6869280

```

cluster_7

```

## photo_sharing      music          cooking      beauty      fashion
##      1.2267360      0.5271313      2.8124735      2.5708935      2.6659507

```

cluster_8

##	chatter	photo_sharing	tv_film	music	shopping
##	1.2243274	0.8780005	0.5222805	0.4020824	1.1005971
##	business	art	small_business		
##	0.4458287	0.4056572	0.4204943		

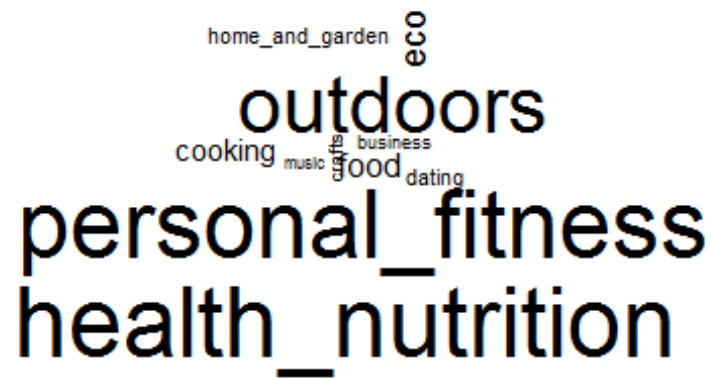
From the above results, we can drop cluster 3 since it has only 1 category and cluster 7 already has similar features.

Now, plotting important features of each cluster in a wordcloud.

```
par(mfrow=c(1,1))
library(wordcloud)

## Loading required package: RColorBrewer

for (i in c(1,2,4,5,6,7,8)) { # skipping cluster 3
wordcloud(colnames(mkt_seg_scale), km_seg$centers[i,], min.freq=0, max.words=
100, scale=c(3,.5))
}
```



A word cloud for the top category. The words are arranged in a roughly triangular shape, with 'personal_fitness' and 'health_nutrition' at the bottom, 'outdoors' in the middle, and 'home_and_garden' at the top. Other words like 'cooking', 'music', 'business', 'good', and 'dating' are scattered in the center.

home_and_garden
outdoors
cooking music business good dating
personal_fitness
health_nutrition



A word cloud for the bottom category. The words are arranged in a roughly triangular shape, with 'automotive' at the top, 'news' in the middle, and 'sports_fandom' at the bottom. Other words like 'politics', 'family', 'tv_film', 'outdoors', 'current_events', 'school', and 'parenting' are scattered in the center.


automotive
politics
family tv_film outdoors current_events school parenting
news
sports_fandom

online_gaming
sports_playing
family art crafts
dating tv_film
home_and_garden
small_business
automotive
college_uni

politics
computers
parenting religion
dating business
eco
small_business
home_and_garden current_events
news food
crafts
sports_playing
travel

religion
parenting
family
music
art
current_events
sports_playing
small_business
crafts
fashion
automotive
computers
eco
beauty
home_and_garden
school
food
sports_fandom

beauty
computers
school
family
art
outdoors
automotive
chatter
home_and_garden
shopping
MUSIC
current_events
eco
dating
sports_playing
business
crafts
photo_sharing
small_business
fashion
cooking



From the clusters obtained, the market segments obtained are the following:

Cluster 1 - Health conscious users

Cluster 2 - Users with more (stereotypical) masculine interests

Cluster 3 - Youngsters (Cluster 3 was dropped and numbers of all others were changed accordingly)

Cluster 4 - Businessmen/Business women

Cluster 5 - Family oriented users

Cluster 6 - Users with more (stereotypical) feminine interests

Cluster 7 - Miscellaneous

These market segments are valuable to NutrientH2O because they now have a better understanding of their customer base by getting a fair idea of what age groups their customers are in, what phase of life they are going through and their hobbies/interests. They can tune their messaging strategy to have customized messages and promotions going out to people based on these interests.